

Stock Prediction:

```
In [ ] :  
  
In [1]:  
#importing  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
%matplotlib inline  
  
In [2]: df = pd.read_csv(r"C:\Users\HP\Downloads\stock_market.csv")  
  
In [3]: df.head()  
  
Out[3]:  
      Date      Open      High      Low      Close  Adj Close  Volume  
0  2018-02-05  262.00000  267.899994  250.029999  254.259995  254.259995  11896100  
1  2018-02-06  247.899997  266.700012  245.000000  265.720001  265.720001  12595800  
2  2018-02-07  266.579987  272.450012  264.329987  264.559998  264.559998  8981500  
3  2018-02-08  267.079987  267.619995  250.000000  250.100006  250.100006  9306700  
4  2018-02-09  253.850006  255.800003  236.110001  249.470001  249.470001  16906900  
  
In [4]: df.tail()  
  
Out[4]:  
      Date      Open      High      Low      Close  Adj Close  Volume  
1004 2022-01-31  401.970001  427.700012  398.200012  427.140015  427.140015  20475600  
1005 2022-02-01  432.959991  458.480011  425.540009  457.130005  457.130005  22542300  
1006 2022-02-02  448.250000  451.980011  426.480011  429.480011  429.480011  14346000  
1007 2022-02-03  421.440002  429.260010  404.279999  405.600006  405.600006  9905200  
1008 2022-02-04  407.309998  412.769989  396.640015  410.170013  410.170013  7782400  
  
In [5]: df.shape  
  
Out[5]: (1889, 7)  
  
In [6]: df.info()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1889 entries, 0 to 1888  
Data columns (total 7 columns):  
 #   Column      Non-Null Count  Dtype    
--  --  --  --  --  --  --  --  
 0   Date       1889 non-null   object   
 1   Open       1889 non-null   float64  
 2   High       1889 non-null   float64  
 3   Low        1889 non-null   float64  
 4   Close      1889 non-null   float64  
 5   Adj Close  1889 non-null   float64  
 6   Volume     1889 non-null   int64    
dtypes: float64(5), int64(1), object(1)  
memory usage: 55.3+ KB  
  
In [7]: df.describe()  
  
Out[7]:  
      Open      High      Low      Close  Adj Close  Volume  
count  1009.000000  1009.000000  1009.000000  1009.000000  1009.000000  1.009000e+03  
mean    419.058673   425.320703   412.374044   419.000733   419.000733  7.570685e+06  
std     108.537532   109.282260   107.555987   108.288999   108.288999  5.465535e+06  
min     233.819998   250.649994   231.229996   233.880005   233.880005  1.144000e+06  
25%     331.489999   336.299989   326.000000   331.619995   331.619995  4.091800e+06  
50%     377.769999   383.010010   370.880005   376.670013   376.670013  5.934600e+06  
75%     509.130005   515.630005   502.529999   509.079987   509.079987  9.322400e+06  
max     692.349976   700.999990   686.090027   691.690002   691.690002  5.890430e+07  
  
In [8]: df.corr()  
  
C:\Users\HP\AppData\Local\Temp\ipykernel_18840\11347222485.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.  
df.corr()  
  
Out[8]:  
      Open      High      Low      Close  Adj Close  Volume  
Open      1.000000  0.998605  0.998508  0.996812  0.996812 -0.415838  
High      0.998605  1.000000  0.998203  0.996551  0.996551 -0.400999  
Low        0.998508  0.998203  1.000000  0.996544  0.996544 -0.421116  
Close      0.996812  0.998551  0.998544  1.000000  1.000000 -0.413362  
Adj Close  0.996812  0.998551  0.998544  1.000000  1.000000 -0.413362  
Volume     -0.415838 -0.400699 -0.422116 -0.413362 -0.413362  1.000000  
  
In [9]: df.nunique()  
  
Out[9]:  
Date      1889  
Open      976  
High      983  
Low        989  
Close      988  
Adj Close  988  
Volume     1885  
dtype: int64  
  
In [10]: df.isnull().sum()  
  
Out[10]:  
Date      0  
Open      0  
High      0  
Low        0  
Close      0  
Adj Close  0  
Volume     0  
dtype: int64  
  
In [11]: df.duplicated().sum()  
  
Out[11]: 0  
  
In [12]: df.columns  
  
Out[12]: Index(['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume'], dtype='object')  
  
In [ ] :  
  
In [ ] :  
  

```

Data Visualization:

```
In [13]: plt.figure(figsize=(30,10))  
sns.lineplot(x="Date",y="Close",data=df)  
plt.grid(axis="y")  
  
In [14]: sns.pairplot(df)  
  
Out[14]: <seaborn.axisgrid.PairGrid at 0x1220b546cd8>  
  
In [15]: sns.scatterplot(df)  
  
Out[15]: <Axes: >  
  
In [16]: sns.jointplot(df)  
  
Out[16]: <seaborn.axisgrid.JointGrid at 0x12208c73d70>  
  
In [17]: sns.heatmap(df.corr(),annot=True)  
  
C:\Users\HP\AppData\Local\Temp\ipykernel_18840\4277794485.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.  
sns.heatmap(df.corr(),annot=True)  
<Axes: >  
  
In [ ] :  
  
In [ ] :  
  

```

Data Modification:

```
In [18]: df  
  
Out[18]:  
      Date      Open      High      Low      Close  Adj Close  Volume  
0  2018-02-05  262.00000  267.899994  250.029999  254.259995  254.259995  11896100  
1  2018-02-06  247.899997  266.700012  245.000000  265.720001  265.720001  12595800  
2  2018-02-07  266.579987  272.450012  264.329987  264.559998  264.559998  8981500  
3  2018-02-08  267.079987  267.619995  250.000000  250.100006  250.100006  9306700  
4  2018-02-09  253.850006  255.800003  236.110001  249.470001  249.470001  16906900  
...      ...      ...      ...      ...      ...      ...      ...  
1004 2022-01-31  401.970001  427.700012  398.200012  427.140015  427.140015  20475600  
1005 2022-02-01  432.959991  458.480011  425.540009  457.130005  457.130005  22542300  
1006 2022-02-02  448.250000  451.980011  426.480011  429.480011  429.480011  14346000  
1007 2022-02-03  421.440002  429.260010  404.279999  405.600006  405.600006  9905200  
1008 2022-02-04  407.309998  412.769989  396.640015  410.170013  410.170013  7782400  
  
1009 rows × 7 columns  
  
In [19]: df.columns  
  
Out[19]: Index(['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume'], dtype='object')  
  
In [20]: x = df.drop(['Date','High','Low','Close','Adj Close','Volume'],axis=1)  
x  
  
Out[20]:  
      Open  
0  262.000000  
1  247.899997  
2  266.579987  
3  267.079987  
4  253.850006  
...      ...  
1004 401.970001  
1005 432.959991  
1006 448.250000  
1007 421.440002  
1008 407.309998  
  
1009 rows × 1 columns  
  
In [21]: y = df['Close']  
y  
  
Out[21]:  
0      254.259995  
1      265.720001  
2      264.559998  
3      250.100006  
4      249.470001  
...      ...  
1884      427.140015  
1885      457.130005  
1886      429.480011  
1887      405.600006  
1888      410.170013  
Name: Close, Length: 1889, dtype: float64  
  
In [ ] :  
  
In [ ] :  
  

```

Machine Learning:

```
In [22]: from sklearn.linear_model  
from sklearn.model_selection import train_test_split  
  
In [23]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)  
  
In [24]: x_train  
  
Out[24]:  
      Open  
643 488.190002  
294 365.109985  
928 633.020020  
131 639.800015  
234 314.570007  
...      ...  
185 305.260010  
74 334.049988  
1006 448.250000  
907 598.570007  
281 358.910004  
  
706 rows × 1 columns  
  
In [25]: x_test  
  
Out[25]:  
      Open  
741 507.350006  
624 492.250000  
455 308.829987  
866 553.969971  
422 270.019989  
...      ...  
290 366.399994  
658 499.989990  
889 515.469971  
513 386.559998  
909 578.169985  
  
303 rows × 1 columns  
  
In [26]: y_train  
  
Out[26]:  
643 490.579987  
294 361.410004  
928 624.540002  
131 641.399998  
234 324.660004  
...      ...  
185 284.839996  
74 331.619995  
1006 429.480011  
907 589.289978  
281 375.220001  
Name: Close, Length: 706, dtype: float64  
  
In [27]: y_test  
  
Out[27]:  
741 509.859985  
624 484.480011  
455 315.549988  
866 542.950012  
422 267.529999  
...      ...  
280 358.779999  
658 483.859985  
889 518.989973  
513 386.000000  
909 582.869995  
Name: Close, Length: 303, dtype: float64  
  
In [28]: print(df.shape)  
print(x_train.shape)  
print(x_test.shape)  
  
(1889, 7)  
(706, 1)  
(303, 1)  
  
In [29]: print(df.shape)  
print(y_train.shape)  
print(y_test.shape)  
  
(1889, 7)  
(706, 1)  
(303, 1)  
  
In [30]: model = linear_model.LinearRegression()  
  
In [31]: model.fit(x_train,y_train)  
  
Out[31]: LinearRegression  
LinearRegression()  
  
In [32]: model.score(x_test,y_test)  
  
Out[32]: 0.9926768195857089  
  
In [ ] :  
  

```

Prediction:

```
In [36]: a = model.predict([[262.000]]) # input is Open and our prediction is close  
print("Stock close:",a)  
Stock close: [262.84197979]  
C:\Users\HP\AppData\Local\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(  
  
In [ ] :  
  
In [ ] :  
  

```